

DL LAB6 REPORT

1. Introduction

在 CV 的 Condition-to-Image 任務中, Diffusion Probabilistic Models 簡稱 DDPM 是用來解這個任務的一種模型。其基本原理是通過一個 forward noising process, 將真實圖像逐步添加 random noise, 直到變成純 noise, 然後再通過 reverse denoising process 根據學習到的參數, 逐步去除噪聲, 最終恢復出高品質的圖像。這兩階段的過程透過馬可夫鏈來逼近資料分佈, 這樣才能夠生成很相似的 output。相比於更老的 GAN, DDPM 在訓練上比較穩定, 且不易發生 mode collapse。此外, DDPM 在 image generation, image inpainting, super-resolution 等任務中表現也比較好, 特別適合需要高質量且 diverse 的圖像輸出, 雖然 sampling speed 相對較慢。在這次的 lab 中, 我們給定條件 (如 “red sphere”, “yellow cube”, “gray cylinder”), 實做了完整的 DDPM pipeline, 從 noise scheduler、模型架構設計到 sampling algorithm。我學到了如何調整 noise schedule、選擇 appropriate loss functions, 收穫很多。

2. Implementation details

DDPM : Conditional DDPM class 中我有使用 Diffusers 套件的 DDPM Scheduler 負責添加 forward noise 和 backward denoise, 用一個基於 Unet model 當做 noise predictor; 再加上 evaluation_model 做驗證; 也有自己寫了 DataLoader 去讀取 iclevrDataset 的 data。Optimizer 採用 Adam, 並配合一個 cosine warmup learning-rate scheduler。在 train() 裡, 對每個 batch 先用 scheduler add noise、再讓 U-Net predict noise, 計算 MSE loss, back propagate 後更新參數, 同時調整 learning rate, 並且在每個 epoch 結束時計算一次 eval acc, 根據表現儲存最佳模型, 最後把訓練過程 (loss、accuracy) 和參數都寫入 record.txt 以便後續分析。程式碼如下圖所示。

```
1 class ConditionalDDPM():
2     def __init__(self, args):
3         self.args = args
4         self.device = args.device
5         self.epoch = args.epoch
6         self.lr = args.lr
7         self.batch_size = args.batch_size
8         self.num_train_timestamps = args.num_train_timestamps
9         self.save_root = args.save_root
10        self.label_embedding_size = args.label_embedding_size
11        self.noise_scheduler = DDPMScheduler(num_train_timestamps=self.num_train_timestamps, beta_schedule="squaredcos_cap_v2")
12        self.noise_predictor = Unet(labels_num=24, embedding_label_size=self.label_embedding_size).to(self.device)
13        self.eval_model = evaluation_model()
14        self.train_dataset = iclevrDataset(root="../dataset/iclevr", mode="train")
15        self.train_loader = DataLoader(self.train_dataset, batch_size=self.batch_size, shuffle=True)
16        self.optimizer = torch.optim.Adam(self.noise_predictor.parameters(), lr=self.lr)
17        self.lr_scheduler = get_cosine_schedule_with_warmup(
18            optimizer=self.optimizer,
19            num_warmup_steps=args.lr_warmup_steps,
20            num_training_steps=len(self.train_loader) * self.epoch,
21            num_cycles=50
22        )
23
```

```

24     def train(self):
25         # print(f"train dataset length: {len(self.train_loader)}")
26         loss_criterion = nn.MSELoss()
27         # training
28         train_loss = []
29         test_acc = []
30         bestacc = 0
31         for epoch in range(1, self.epoch+1):
32             epoch_loss = 0
33             for x, y in tqdm(self.train_loader):
34                 x = x.to(self.device)
35                 y = y.to(self.device)
36                 noise = torch.randn_like(x)
37                 timestamp = torch.randint(0, self.num_train_timestamps - 1, (x.shape[0], ), device=self.device).long()
38                 noise_x = self.noise_scheduler.add_noise(x, noise, timestamp)
39                 perdn_noise = self.noise_predictor(noise_x, timestamp, y)
40                 loss = loss_criterion(perdn_noise, noise)
41                 loss.backward()
42                 nn.utils.clip_grad_value_(self.noise_predictor.parameters(), 1.0)
43                 self.optimizer.step()
44                 self.lr_scheduler.step()
45                 self.optimizer.zero_grad()
46                 self.lr = self.lr_scheduler.get_last_lr()[0]
47                 epoch_loss += loss.item()
48             epoch_loss /= len(self.train_loader)
49             train_loss.append(epoch_loss)
50             print('Epoch ' + str(epoch) + ' Loss: ' + str(epoch_loss))
51             # acc = self.evaluate(epoch, testing_dataset="test")
52             acc = self.evaluate( testing_dataset="test")
53             test_acc.append(acc)
54             print('Epoch ' + str(epoch) + ' Accuracy: ' + str(acc))
55             if acc > bestacc or acc > 0.8 :
56                 if acc > bestacc :
57                     bestacc = acc
58                     self.save(os.path.join(self.args.ckpt_root, f"epoch={epoch}.ckpt"), epoch)
59             f = open(os.path.join(self.args.ckpt_root, f"record.txt"), 'w')
60             print('Arguments:', file=f)
61             for arg, value in vars(self.args).items():
62                 print(f"[arg]: {value}", file=f)
63             print(file=f)
64
65             print('Loss', file=f)
66             for i in range(self.epoch):
67                 print(train_loss[i], ',', end=' ', file=f)
68                 if i % 10 == 9:
69                     print(file=f)
70             print(file=f)
71
72             print('Accuracy', file=f)
73             for i in range(self.epoch):
74                 print(test_acc[i], ',', end=' ', file=f)
75                 if i % 10 == 9:
76                     print(file=f)
77             print(file=f)
78
79             f.close()

```

noise predictor 的部分我寫在[UNET.py](#)中，如下圖。

程式碼中有使用diffuser的UNet2DModel。Unet的input 有RGB還有condition，所以對condition進行shapeexpand 來align input image shape，之後再將他們concatenate，因此input channel會是3(RGB)+24(Condition)。timeembedding 則直接將 timestamp 輸入到Unet2DModel，model 便

會自動處理。我實作noise_predicter的程式碼如下頁 圖所示。


```
1 class Unet(nn.Module):
2     def __init__(self, labels_num = 24, embedding_label_size=4) -> None:
3         super().__init__()
4         self.label_embedding = nn.Embedding(labels_num, embedding_label_size)
5         self.model = UNet2DModel(
6             sample_size = 64,
7             # in_channels = 3 + labels_num * embedding_label_size,
8             in_channels = 3 + labels_num,
9             out_channels = 3,
10            time_embedding_type = "positional",
11            layers_per_block = 2,
12            block_out_channels = (128, 128, 256, 256, 512, 512), # number of output channels for each UNet block
13            down_block_types = (
14                "DownBlock2D", # a regular ResNet downsampling block
15                "DownBlock2D",
16                "DownBlock2D",
17                "DownBlock2D",
18                "AttnDownBlock2D", # a ResNet downsampling block with spatial self-attention
19                "DownBlock2D",
20            ),
21            up_block_types = (
22                "UpBlock2D", # a regular ResNet upsampling block
23                "AttnUpBlock2D", # a ResNet upsampling block with spatial self-attention
24                "UpBlock2D",
25                "UpBlock2D",
26                "UpBlock2D",
27                "UpBlock2D",
28            ),
29        )
30    def forward(self, x, t, label):
31        bs, c, w, h = x.shape
32        embeded_label = label.view(bs, label.shape[1], 1, 1).expand(bs, label.shape[1], w, h)
33        unet_input = torch.cat((x, embeded_label), 1)
34        unet_output = self.model(unet_input, t).sample
35        return unet_output
```

3. Results and discussion

- Show your synthetic image grids (total 16%: 8% * 2 testing data)
 - test.json



- new_test.json



- denoising process image with the label set ["red sphere", "cyan cylinder", "cyan cube"]



- Discussion of your extra implementations or experiments (10%)



有測試了兩組hyperparameters，其結果如下面四張圖可以看到，

基本上epoch 150 已經能夠達到還不錯的accuracy，只是再繼續train到300其實大約就能

從Acc 0.7左右提升至Acc 0.8 左右，可以觀察到DDPM其實相較於GAN很好train，不會因為hyper parameters 的一點點差異就導致model collapse。

lr	1e-4	2e-4
epoch	150	300
acc curve	<p>Accuracy Curve</p> <p>Y-axis: Accuracy (0.1 to 0.8)</p> <p>X-axis: Epoch (0 to 150)</p> <p>Legend: Accuracy</p>	<p>Accuracy Curve</p> <p>Y-axis: Accuracy (0.1 to 0.9)</p> <p>X-axis: Epoch (0 to 300)</p> <p>Legend: Accuracy</p>

loss curve



4.Experimental results

test.json acc : 0.819

new_test.json acc : 0.905

```
(base) winston@gpu9:/project2/winston/lab/lab6/DDPM$ python main.py --test True --ckpt_path ./epoch=220.ckpt --ckpt_root ./save_model/2
In testing mode : True
testing_dataset : test
1000it [00:37, 26.99it/s]
DDPM test.json Accuracy : 0.819
testing_dataset : new_test
1000it [00:36, 27.36it/s]
DDPM new_test.json Accuracy : 0.905
1000it [00:10, 94.39it/s]
```